Prediction accuracy and candidate recovery are different goals

30 MoleculeACE target–endpoint tasks · 29 target clusters · 9,802 held-out compound records

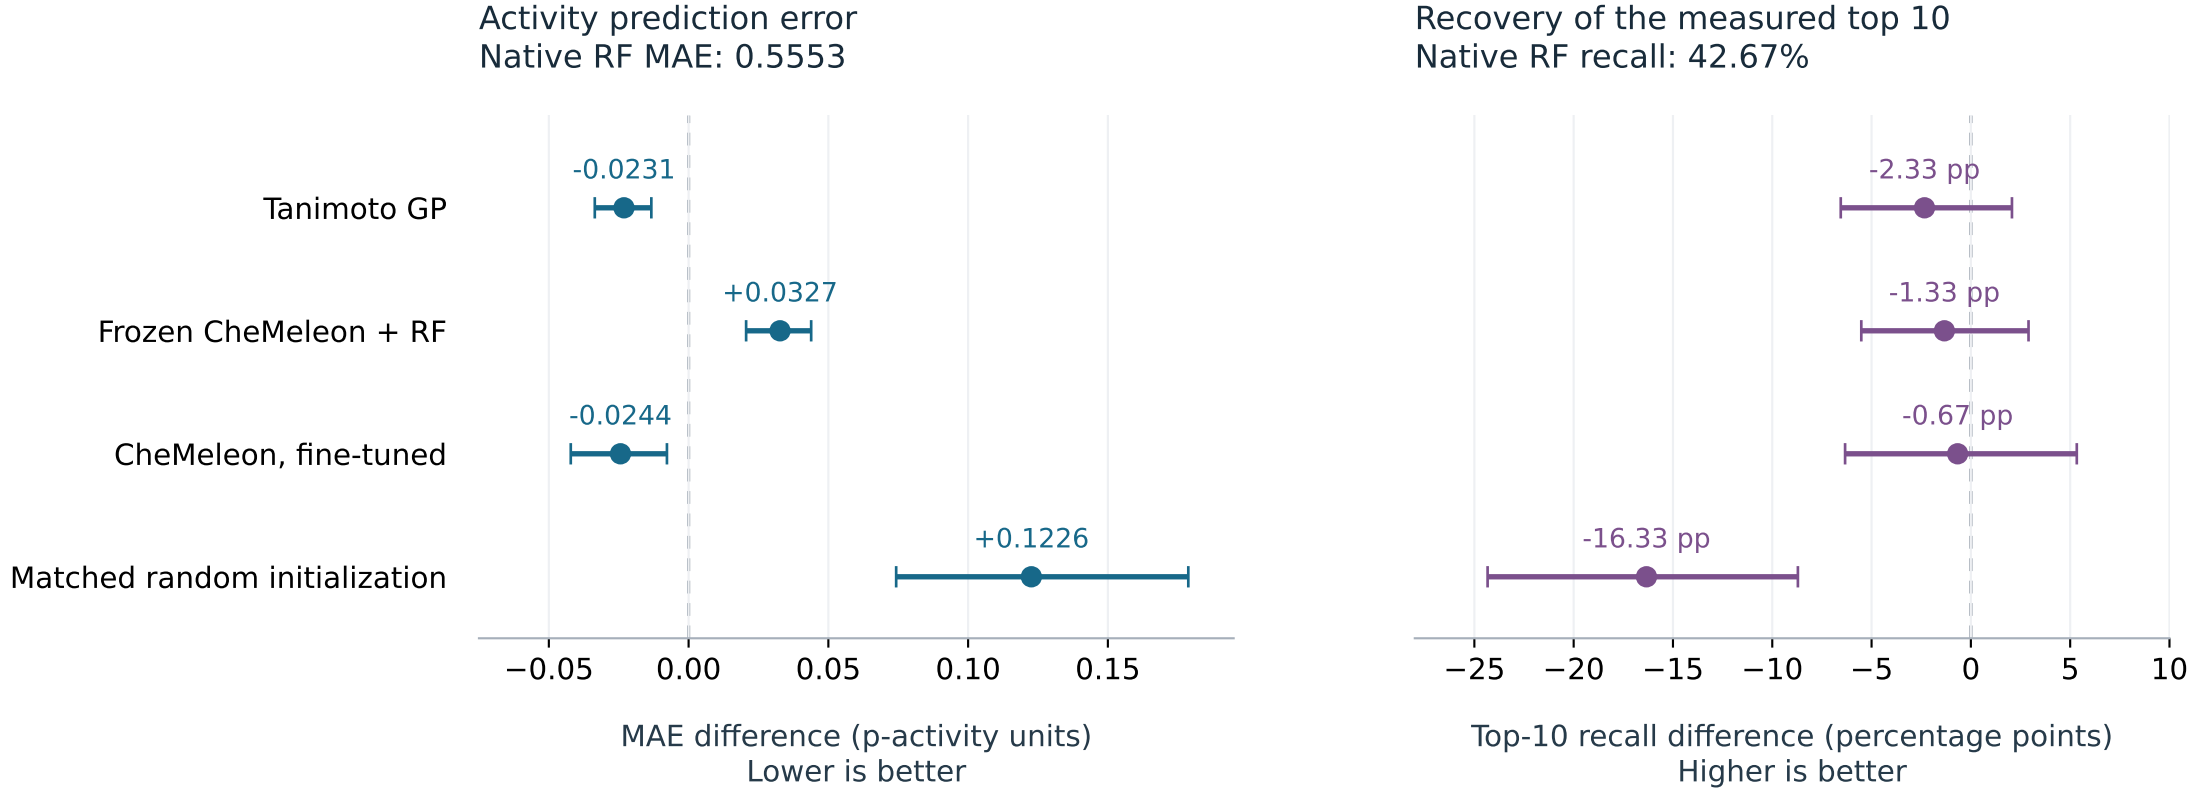

Points: paired task-average differences. Bars: 95% target-cluster bootstrap intervals (5,000 draws).

Development experiments tried sequentially on the same benchmark; one fit per task. All five model variants are retained.